

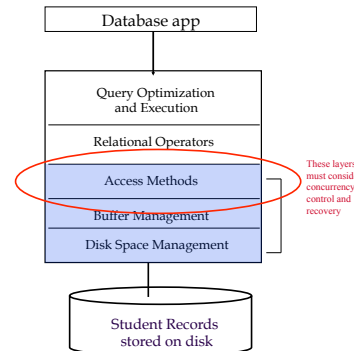
CS186: Introduction to Database Systems

Lecture 3 – File Management (Book Ch: 9.5-7)

Michael Franklin
Fall 2013



Context



Files of Records

- Disk blocks are the interface for I/O, but...
- Higher levels of DBMS operate on *records*, and *files of records*.
- **FILE**: A collection of **pages**, each containing a number of records. The File API must support:
 - insert/delete/modify** record
 - fetch** a particular record (specified by *record id*)
 - scan** all records (possibly with some conditions on the records to be retrieved)
- Typically: **file page size = disk block size = buffer frame size**

“MetaData” - System Catalogs

- How to impose structure on all those bytes??
 - MetaData: “Data about Data”
 - For each relation:
 - name, file location, file structure (e.g., Heap file)
 - attribute name and type, for each attribute
 - index name, for each index
 - integrity constraints
 - For each index:
 - structure (e.g., B+ tree) and search key fields
 - For each view: view name and definition
 - Plus statistics, authorization, buffer pool size, etc.
- ❖ *Q: But how to store the catalogs???*

Catalogs are Stored as Relations!


attr_name	rel_name	type	position
attr_name	Attribute_Cat	string	1
rel_name	Attribute_Cat	string	2
type	Attribute_Cat	string	3
position	Attribute_Cat	integer	4
sid	Students	string	1
name	Students	string	2
login	Students	string	3
age	Students	integer	4
gpa	Students	real	5
fid	Faculty	string	1
frame	Faculty	string	2
sal	Faculty	real	3

Attr_Cat(attr_name, rel_name, type, position, length)

It's a bit more complicated...

```
Terminal -- psql -- 99x28
josh# \dt pg_attribute
No matching relations found.
josh# \d pg_attribute
Table "pg_catalog.pg_attribute"
  Column   | Type          | Modifiers
-----+-----+-----
attrelid   | oid           | not null
attname    | name          | not null
atttypid   | oid           | not null
attstattarget | integer      | not null
attlen     | smallint     | not null
attnum     | smallint     | not null
attndims   | integer      | not null
attcacheoff | integer      | not null
atttypmod  | integer      | not null
attbyval   | boolean      | not null
attstorage | "char"       | not null
attalign   | "char"       | not null
attisnull  | boolean      | not null
attisuser  | boolean      | not null
attisdropped | boolean      | not null
attislocal | boolean      | not null
attinhcount | integer      | not null
Indexes:
    "pg_attribute_relid_attname_index" UNIQUE, btree (attrelid, attname)
    "pg_attribute_relid_attnum_index" UNIQUE, btree (attrelid, attnum)
josh#
```

Record Formats: Fixed Length




- Information about field types same for all records in a file; stored in *system catalogs*.
- Finding i 'th field done via arithmetic.


Record Formats: Variable Length



- Two alternative formats (# fields is fixed):



Fields Delimited by Special Symbols



Array of Field Offsets

- Second offers direct access to i 'th field, efficient storage of *nulls* (special *don't know* value); some directory overhead.


How to Identify a Record?



- The Relational Model doesn't expose "pointers", but that doesn't mean that the DBMS doesn't use them internally.
- Q: Can we use *memory addresses* to permanently "point" to records?
- Systems use a "Record ID" or "RecID"

Typically: Record ID = *<page id, slot #>*


Page Formats: Fixed Length Records



In first alternative, free space management requires record movement.

Changes RIds - may not be acceptable.

"Slotted Page" for Variable Length Records




- Slot contains: [offset (from start of page), length]
 - both in bytes
- Record id = *<page id, slot #>*
- Page is full when data space and slot array meet.

Slotted Page (continued)




- When need to allocate:
 - If enough room in free space, use it and update free space pointer.
 - Else, try to compact data area, if successful, use the freed space.
 - Else, tell caller that page is full.
- Advantages:
 - Can move records around in page without changing their record ID
 - Allows lazy space management within the page, with opportunity for clean up later

Slotted page (continued)




- What's the biggest record you can add to the above page without compacting?
 - Need 2 bytes for slot: [offset, length] plus record.

Slotted page (continued)




- What's the biggest record you can add to the above page without compacting?
 - Need 2 bytes for slot: [offset, length] plus record.
- What happens when a record needs to move to a different page?
 - Leave a “tombstone”, pointing to new page & slot.
 - Record id remains unchanged –one hop max.

Slotted page (continued)



- What's the biggest record you can add to the above page **with compacting**?
 - Need 2 bytes for slot: [offset, length] plus record.

Slotted page (continued)



- What do you do if a record needs to move to a different page?
 - Leave a special “tombstone” object in place of record, pointing to new page & slot.
 - Record id remains unchanged
- What if it needs to move again?
 - Update the original tombstone – so one hop max.

So far we've organized:

- Fields into Records (fixed and variable length)
- Records into Pages (fixed and variable length)

Now we need to organize Pages into Files

Alternative File Organizations

Many alternatives exist, *each good for some situations, and not so good in others*:

Heap files: Unordered. Fine for file scan retrieving all records. Easy to maintain.

Sorted Files: Best for retrieval in *search key* order, or if only a 'range' of records is needed. Expensive to maintain.


Clustered Files (with Indexes): A compromise between the above two extremes.

Unordered (Heap) Files




- Simplest file structure contains records **in no particular order**.
- As file grows and shrinks, pages are allocated and de-allocated.
- To support record level operations, we must:
 - keep track of the **pages** in a file
 - keep track of **free space** on **pages**
 - keep track of the **records** on a page
- Can organize as a list, as a directory, a tree, ...

Heap File Implemented as a List



- The Heap file name and header page id must be stored persistently.
The catalog is a good place for this.
- Each page contains 2 'pointers' plus data.

Heap File Using a Page Directory



- The entry for a page can include the number of free bytes on the page.
- The directory is a collection of pages; linked list implementation is just one alternative.
- **Q: How to find a particular record in a Heap file???**

Cost Model for Analysis



- Average-case analysis; based on several simplistic assumptions.
 - Often called a **"back of the envelope"** calculation.
- We ignore CPU costs, for simplicity:
 - B**: The number of data blocks
 - R**: Number of records per block
- We simply count number of disk block I/O's
 - ignores gains of pre-fetching and **sequential access**; thus, even I/O cost is only loosely approximated.

➡ *Good enough to show some overall trends!*

Some Assumptions in the Analysis



- Single record insert and delete.
- Equality selection - exactly one match (what if more or less???)
- For Heap Files we'll assume:
 - Insert always appends to end of file.
 - Delete just leaves free space in the page.
 - Empty pages are not deallocated.

Average Case I/O Counts for Operations ($B = \#$ disk blocks in file)



	Heap File	Sorted File	Clustered File
Scan all records	B		
Equality Search (1 match)	$0.5 B$		
Range Search	B		
Insert	2		
Delete	$0.5B+1$		

Sorted Files



- Heap files are **lazy** on **update** - you end up paying on searches.
- **Sorted files** **eagerly** maintain the file on **update**.
 - The opposite choice in the trade-off
- Let's consider an extreme version
 - No gaps allowed, pages fully packed always
 - Q: How might you relax these assumptions?
- Assumptions for our BotE Analysis:
 - Files compacted after deletions.
 - Searches are on sort key field(s).

Average Case I/O Counts for Operations ($B = \#$ disk blocks in file)



	Heap File	Sorted File	Clustered File
Scan all records	B	B	
Equality Search (1 match)	$0.5 B$	$\log_2 B$ (if on sort key) $0.5 B$ (otherwise)	
Range Search	B	$(\log_2 B) + \text{selectivity} * B$	
Insert	2	$(\log_2 B) + B$	
Delete	$0.5B + 1$	Same cost as Insert	

File Structure Summary



- File Layer manages access to records in pages.
 - Record and page formats depend on fixed vs. variable-length.
 - Free space management is an important issue.
 - **Slotted page format** supports variable length records and allows records to move on page.
- Many alternative file organizations exist, each appropriate in some situation.
 - We looked at Heap and Sorted so far.
 - If selection queries are frequent, sorting the file or building an *index* is important.
- Back of the envelope calculations are imprecise, but can expose fundamental systems tradeoffs.
 - A technique that you should become comfortable with!
- Next up: Indexes.